

# Visualizing Token Importance in Large Language Models

A Review of the State of and a Framework for Understanding Token Attribution in LLMs

**Ruben Kaiser**✉

Software and Systems Engineering, Technical University of Munich

✉ ruben.kaiser@tum.de

April 15, 2025

**Supervisor:** Xingcheng Chen

**Examiner:** Prof. Dr. Andrea Stocco

**Abstract** — Large Language Models have become increasingly integrated into critical domains, yet their internal processes remain largely opaque. While various attribution methods for transparency exist, visualization tools that effectively balance technical faithfulness with interpretability for non-experts are absent from the current landscape. Existing tools such as BertViz, Polyjuice, and Inseq suffer from specific limitations: incompatibility with modern architectures, post-hoc rationalizations instead of faithful explanations, and static rather than interactive visualizations.

This paper conducts a systematic analysis of current token attribution tools through the lens of Human-Computer Interaction (HCI) principles, drawing from comprehensive reviews of explainability methods. Based on identified gaps, it proposes an integrated visualization framework with three technical innovations: (1) a mathematically formalized job-based architecture enabling parallel and efficient attribution processing across multiple methods, (2) a dynamic visualization interface supporting context-sensitive exploration through token hovering, window adjustment, and direct model comparison, and (3) an extensible system designed to simultaneously visualize diverse attribution methods with appropriate visual encodings.

The framework is validated through TokenLense<sup>1</sup>, a proof-of-concept implementation that demonstrates interactive token relationship visualization using background highlighting and connection lines, comparative analysis across models and methods, and performance benchmarking of attribution techniques. Implementation challenges revealed critical insights about model compatibility and workflow integration that inform future directions, including AR visualization techniques to address information density challenges in complex models. The proposed framework addresses the tension between faithfulness and plausibility in LLM interpretability, potentially enabling stakeholders from model developers to domain experts in healthcare, legal, and educational contexts to gain actionable insights into how LLMs process information. Such transparency is essential for identifying reliability issues, mitigating biases, and ensuring proper alignment as these systems become increasingly embedded in critical decision-making processes.

## 1 Introduction

### 1.1 Background

The explosion in performance of Large Language Models (LLMs) has marked a transformative era in artificial intelligence [1–5]. The resulting widespread adoption can be largely attributed to the dramatic increase in model scale, with larger models consistently demonstrating enhanced capabilities across domains: Building upon the Transformer architecture introduced by Vaswani et al. [6], language models have evolved from the initial GPT [7] to sophisticated systems like BERT [8], GPT-3 [9], and more recent iterations such as GPT-4 [10], LLaMA-3 [11], and Claude [12], demonstrating remarkable capabilities across a wide range of tasks.

In **healthcare**, these models assist with medical diagnosis and clinical documentation [13]. **Financial institutions** leverage LLMs for market analysis and risk assessment, with specialized models like BloombergGPT showing enhanced performance in finance-specific tasks [14]. **Educational** applications span personalized

---

<sup>1</sup>TokenLense on GitHub

tutoring and curriculum development (see Abd-Alrazaq et al. for medical education [15]), while **legal professionals** use LLMs for document review and contract analysis [16]. **Creative industries** have embraced these models for content generation across multiple formats [17]. This widespread integration across critical infrastructure and decision-making processes continues to reshape economic systems, information environments, and social dynamics [18].

The variety of task is enabled by the ever increasing sizes of the foundation models and the amount of available training data, and - by correlation [19] - their increase in performance, with consistent scaling laws predicting performance improvements as parameters, training data, and computational resources increase [20].

However, this rapid advance in capabilities has outpaced our understanding of how these systems function. As Liao and Vaughan argue, LLMs present "*looming risks to society itself*," [21] necessitating that "*transparency must be a central pillar*" [21] of responsible development and deployment. Weidinger et al. [18] define six risk areas, spanning from (miss-)information to environmental and social-economic harms.

Unlike simpler machine learning models whose operations can be readily inspected, LLMs operate as complex, distributed systems whose internal representations and decision-making processes remain largely opaque. This opacity is particularly concerning given their above-mentioned growing integration into high-stakes domains where errors or biases could have significant consequences for individuals and communities [22] as well as society as whole.

## 1.2 Challenges in LLM Interpretability

Several key components lead to complications in developing effective interpretability methods for LLMs:

1. The sheer **scale and complexity** of these models—often containing hundreds of billions of parameters—render many traditional explainability approaches computationally *infeasible or inadequate* for capturing the multi-layered transformations that occur throughout the system [23]. A complete analysis of parameter interactions across numerous attention heads and layers quickly becomes infeasible, and even when done, often don't convey meaning or doesn't present the full picture [24].
2. **Competing interpretability paradigms** approach the problem from different perspectives, each with their own methodological assumptions. Some focus on mechanistic interpretability, seeking to reverse-engineer specific circuits within the model [25]. Others adopt a more functional approach, analyzing the relationship between inputs and outputs without attempting to map the exact computational paths [26]. And others again rely on the model itself to explain its "reasoning" [27]. These divergent approaches, while valuable in diversity, have created a fragmented research landscape that can be difficult for practitioners to navigate.
3. The **proprietary nature** of many leading LLMs (OpenAI, Google DeepMind, Anthropic) presents a significant barrier to open research. With limited access to model weights, training data, and internal representations, researchers must often work with incomplete information about the systems they seek to understand. [23]
4. The **contextual and emergent nature** of representations in LLMs poses unique challenges. Unlike traditional feature importance methods that can identify fixed, context-independent relationships between inputs and outputs, LLMs construct dynamic, context-sensitive representations that may vary substantially across different inputs and generation steps [28]. This contextuality makes it difficult to generalize explanations across different use cases or to provide simplified yet accurate accounts of model behavior.

We will look at some specific examples for visualisation attempts in section 2.

## 1.3 Current Approaches to Interpretability

The field of LLM interpretability encompasses a diverse range of methods that attempt to illuminate different aspects of model behavior and knowledge. Zhao et al. [23] and Cambria et al. [22] provide comprehensive surveys of existing approaches, highlighting both their strengths and limitations.

**Feature attribution methods** aim to identify which input tokens most influenced a particular prediction or generation. Techniques such as integrated gradients [29], LIME [30], and SHAP [31] have been adapted for language models, although their computational requirements can become prohibitive for very large models. These methods often struggle to account for the complex interactions between tokens in contextual representations, potentially yielding misleading attributions in certain scenarios.

**Attention visualization tools** offer insights into how models allocate focus across input elements. Implementations like BertViz [32] allow users to inspect attention patterns across different layers and heads, revealing structural regularities in how models process linguistic information. However, a substantial debate exists regarding whether attention weights genuinely reflect causal relationships in model reasoning [24,33,34].

**Probing studies** examine whether specific linguistic properties are encoded in model representations, leveraging classifiers trained on model activations to detect the presence of syntactic, semantic, or factual knowledge [35]. While valuable for understanding what information models capture, probing methods typically yield limited insights into how this information is utilized during actual generation or prediction tasks.

**Concept-based approaches** investigate how high-level concepts are represented across model components [36,37]. These methods bridge the gap between low-level activations and human-interpretable ideas, offering a more abstract but potentially more useful characterization of model knowledge. However, defining and validating appropriate concepts remains challenging, particularly for the diverse and contextual representations in LLMs.

Despite this array of techniques, measuring and visualizing token importance in LLMs presents persistent difficulties. As models grow in size and complexity, traditional attribution methods become computationally expensive and may fail to capture the hierarchical, distributed nature of token representations [23]. Furthermore, token importance can vary dramatically with context, making static visualizations potentially misleading without appropriate qualification.

## 1.4 Problem Statement

Given these challenges and current limitations ( $\rightarrow$  section 2) in interpretability and the potential social implications of widespread LLM use, there is an urgent need for more effective and accessible tools to visualize and analyze token importance in LLMs. This work addresses this gap by proposing a roadmap and requirements for a visualization tool that builds upon existing research but reimagines interpretability through human-computer interaction principles [38,39].

This approach aims to:

1. Make complex model behavior accessible through intuitive interfaces that don't require deep technical expertise
2. Generate actionable insights that can inform responsible development and deployment
3. Bridge the gap between technical precision and practical usability

By democratizing access to LLM interpretability tools, we can enable a wider range of stakeholders—from domain experts in healthcare and education to regulators and end users—to meaningfully engage with questions of model behavior, bias, and limitations as these systems continue to transform our information environment and social structures.

## 2 Current State of Token Importance Visualization

*We do not introduce the concepts of tokens in this paper, but the reader is invited to read either Sennrich et al. [40] for their work on subword tokenization for neural machine translation, or Vaswani et al. [6] for their explanation of tokens in the context of transformer architectures.*

## 2.1 Current Approaches to Token Importance Visualization

The rapidly expanding domain-agnostic involvement of Large Language Models (LLMs) necessitates robust interpretability methods [21]. In this section, we analyze current approaches to token visualization, focusing on their theoretical foundations and practical implementations.

As foundation to understanding and evaluating existing tools, we must first examine the attribution methods they employ to quantify token importance.

### 2.1.1 Token Attribution Methods

Token attribution in LLMs aims to quantify how individual input tokens (words, subwords, or characters) contribute to a model’s output. Several methodologies have emerged, of which the most prominent ones are:

**Integrated Gradients** Introduced by Sundararajan et al. [29], Integrated Gradients (IG) addresses limitations in standard gradient-based attribution by satisfying important axioms: sensitivity and implementation invariance. IG calculates the path integral of gradients from a baseline input to the actual input:

$$\text{IntegratedGrads}_i(x) = (x_i - x'_i) \times \int_{\alpha=0}^1 \frac{\partial F(x' + \alpha \times (x - x'))}{\partial x_i} d\alpha \quad (1)$$

IG’s computational requirements are substantial for LLMs with billions of parameters, as it requires multiple forward and backward passes to approximate the integral [23].

**SHAP** SHapley Additive exPlanations (SHAP) [31] unifies various attribution methods using concepts from game theory: It treats features as players in a cooperative game and assigns importance based on their marginal contribution to predictions. While SHAP provides desirable properties like local accuracy and consistency, its exact calculation has exponential complexity.

**Attention mechanisms** Transformer-based LLMs rely on self-attention mechanisms that assign weights to tokens during processing [6]. These weights present an intuitive target for interpretation, visualized in tools like BertViz [32]. However, significant debate exists regarding attention’s reliability for explanation [24].

**Perturbation and surrogate methods** Perturbation-based methods assess token importance by measuring how removing, masking, or altering tokens affects the model’s output. Surrogate methods like LIME [30] create locally faithful explanations by approximating complex models with simpler, interpretable ones.

### 2.1.2 Explainability Principles for Visualization Tools

Effective LLM visualization tools must balance technical rigor with usability. Two critical dimensions for evaluating explanations are according to Zhao et al. [23]:

- **Faithfulness:** How accurately an explanation reflects the model’s actual decision process.
- **Plausibility:** How understandable and convincing an explanation appears to humans.

These dimensions often exist in tension: highly plausible explanations may not faithfully represent the model’s internal processes. A tangible example is Chain of Thought (CoT) reasoning, which is very natural to understand but can generate plausible yet unfaithful explanations [27].

From an HCI perspective, Liao et al. [21] emphasize human-centered transparency approaches, including: progressive disclosure of complexity; appropriate visual encodings; interactive exploration capabilities; and support for developing accurate mental models over time.

## 2.2 Existing Tools and Approaches

The field of LLM interpretability has seen rapid development of specialized tools and frameworks for token attribution visualization. Table 1 summarizes key visualization tools, their approaches, and limitations.

**Table 1** Comparison of Token Importance Visualization Tools

| Tool                           | Key Features & Attribution                                                                                                                                                                                                                         | Strengths                                                                                                                                                                                                                  | Limitations                                                                                                                                                                                                                 |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>BertViz</b> [32]            | Multi-scale visualization: attention-head view, model view, neuron view; Birds-eye view of attention across layers and heads<br>( <i>Attention weights</i> )                                                                                       | Detailed visualization of attention patterns across model components; Examines individual components within attention mechanism                                                                                            | Poor correlation between attention and feature importance; Models can maintain identical predictions with drastically altered attention patterns; Limited applicability to current SOTA models                              |
| <b>Polyjuice</b> [41]          | Counterfactual example generation; Fine-tuned on diverse datasets; Produces contextually appropriate counterfactuals<br>( <i>Contrastive analysis</i> )                                                                                            | Illuminates model behavior through contrastive analysis; Diverse, contextually appropriate counterfactual examples                                                                                                         | Based on GPT-2 (architectural constraints); Computationally expensive for complex prompts; Lacks integration with other attribution methods; Cannot readily generalize to modern architectures                              |
| <b>LMExplainer</b> [42]        | Knowledge graph integration; Graph attention networks (GAT); Two-stage explanation generation; Debugging framework for evaluation quality<br>( <i>Knowledge-grounded attention</i> )                                                               | Structured knowledge-grounded explanations; Addresses both why options were chosen and others rejected; Performs well on CommonsenseQA and OpenBookQA datasets                                                             | Uncertain relationship between extracted reason-elements and actual decision processes; Limited evaluation scale; Knowledge graph pruning and parameter settings lack specification                                         |
| <b>Chain-of-Thought</b> [27]   | Natural language step-by-step reasoning; Self-reported explanation process<br>( <i>Self-reported explanation</i> )                                                                                                                                 | Human-interpretable reasoning format; Provides tangible explanations                                                                                                                                                       | Generates plausible yet unfaithful explanations; Performance dropped up to 36% on tasks with biased prompts; Models justify stereotype-aligned answers without mentioning stereotypes                                       |
| <b>Inseq</b> [43]              | Centralized access to multiple attribution methods; Model-architecture independence; Structured output formats (JSON); Standard static visualizations<br>( <i>IG, SHAP, LIME, Attention</i> )                                                      | Works with all major Hugging Face models; Enables comparison across attribution techniques; Architecture-independent design; Built on PyTorch and transformers library                                                     | Limited to static visualizations (heatmaps); Lacks interactive exploration interfaces; Provides primarily static heatmaps rather than dynamic exploration tools                                                             |
| <b>CodeQ</b> [44]              | Rationalization of code predictions (minimal token subsets); Mapping to programming concepts via AST; Reduction via interpretability tensors; Domain-specific constructs<br>( <i>Domain-specific attribution</i> )                                 | Domain-specific human-understandable explanations; High perceived usefulness for fine-tuning (89%) and output interpretation (84%); Detects "overinterpretation" phenomena                                                 | Poor alignment with human reasoning (Jaccard: 0.074); Only 52-59% found visualizations easy to interpret; Computational expense; Feature splitting across concepts                                                          |
| <b>Attribution Graphs</b> [45] | Cross-layer transcoders replace neurons with interpretable features; Local replacement models with frozen attention; Attribution graphs trace causal paths; Graph pruning; Interactive visualization interface<br>( <i>Causal graph analysis</i> ) | Traces causal paths from inputs through features to outputs; Supports testing mechanistic hypotheses; Identified mechanisms behind factual recall, arithmetic, and reasoning; Allows researchers to group related features | Cannot explain attention formation (query-key interactions); Substantial computational cost (thousands of GPU hours); Compounding errors for longer-path mechanisms; "Dark matter" in model computation remains unexplained |

## 2.3 Limitations of Current Approaches

Current token importance visualization approaches face several key limitations:

- **Faithfulness-plausibility tension:** Tools often sacrifice faithful representation of model internals for human-understandable explanations. Recent work by Palacio et al. [44] quantifies this tension, revealing significant misalignment between model rationales and human reasoning patterns.

- **Computational and mechanism constraints:** Methods like Integrated Gradients and SHAP face substantial computational barriers with billion-parameter models. Attribution graphs [45] identify blind spots in explaining attention patterns, while others miss inhibitory relationships between features.
- **Interactivity and ecosystem fragmentation:** Most current tools provide static visualizations rather than supporting interactive exploration. The fragmented ecosystem (different tools for different aspects) creates integration challenges for researchers conducting thorough analyses.
- **Abstraction and domain integration challenges:** Finding appropriate levels of abstraction remains difficult, with recent methods identifying feature splitting where conceptually unified ideas are distributed across multiple features. Generic approaches rarely connect token importance to broader semantic structures or domain-specific knowledge.

These limitations highlight the need for integrated approaches that provide both global overviews and detailed, interactive exploration capabilities while maintaining fidelity to the model’s actual processing mechanisms. Future visualization tools should aim to bridge the gap between faithful representation of model mechanisms and human-interpretable explanations, while providing domain-appropriate abstractions that connect to existing knowledge structures.

### 3 Core Mathematical Concepts

Based on the identified challenges in LLM interpretability and limitations of existing approaches, this section formalizes the requirements and development goals for the proposed attribution visualization framework.

#### 3.1 Theoretical Foundation Requirements

Let  $\mathcal{M} = \{m_1, m_2, \dots, m_n\}$  be the set of available (transformer based) generation models. Let  $\mathcal{F} = \{f_1, f_2, \dots, f_k\}$  be the set of attribution methods, encompassing gradient-based, internals-based, and perturbation-based approaches.

For any attribution function  $\phi$ , the mapping of an input-output pair to a matrix of importance scores is formalized as:

$$\phi_f^m(X, Y, o) = \mathbf{A}^{m,f,X,Y} \in \mathbb{R}^{n \times t} \quad (2)$$

where:

- $X = (x_1, x_2, \dots, x_n)$  represents the input token sequence
- $Y = (y_1, y_2, \dots, y_t)$  represents the output token sequence
- $\mathbf{A}_{i,j}^{m,f,X,Y}$  denotes the attribution score of input token  $x_i$  for output token  $y_j$
- $o \in \mathcal{O}$  represents the set of options applied to attribution calculations

Given the limitations of individual attribution methods highlighted in Section 2, the system must support multiple attribution approaches within a unified interface. Cross-model compatibility must be ensured through abstraction of model-specific details, handled by Inseq (see Appendix A.6).

#### 3.2 Computational and Processing Requirements

Attribution calculations for LLMs are computationally intensive, necessitating a job-based processing architecture. A job  $J$  is defined as:

$$J = (\mathcal{M}', \mathcal{F}', i, o) \quad (3)$$

where:

- $\mathcal{M}' \subseteq \mathcal{M}$  is the subset of selected models

- $\mathcal{F}' \subseteq \mathcal{F}$  is the subset of selected attribution methods
- $i \in \mathcal{I}$  is a single input text, where  $\mathcal{I}$  represents the set of possible input texts
- $o \in \mathcal{O}$  is a set of options applied to all model-method combinations

Each job processes all combinations in the Cartesian product  $\mathcal{M}' \times \mathcal{F}'$ . The system must support asynchronous processing, parallel execution, effective resource management, and result caching for improved user experience. Detailed API specifications are provided in Appendix A.1.

### 3.3 Aggregation Functions

To enable meaningful interactive analysis, several key aggregation functions are defined:

- **Input-centered aggregation** (when focusing on input token  $x_k$ ):

$$\mathbf{A}_{x_k, c}(j) = \sum_{i=\max(1, k-\lfloor c/2 \rfloor)}^{\min(n, k+\lfloor c/2 \rfloor)} \mathbf{A}_{i, j} \quad (4)$$

- **Output-centered aggregation** (when focusing on output token  $y_k$ ):

$$\mathbf{A}_{y_k, c}(i) = \sum_{j=\max(1, k-\lfloor c/2 \rfloor)}^{\min(t, k+\lfloor c/2 \rfloor)} \mathbf{A}_{i, j} \quad (5)$$

- **Token importance aggregation** (for default visualization):

$$I(x_i) = \sum_{j=1}^t \mathbf{A}_{i, j} \quad \text{and} \quad I(y_j) = \sum_{i=1}^n \mathbf{A}_{i, j} \quad (6)$$

Additional aggregation functions for multi-dimensional methods and multi-method combinations are detailed in Appendix A.3.

### 3.4 Visualization Requirements

The system must visually encode attribution scores  $\mathbf{A}_{i, j}$  through appropriate perceptual channels while maintaining text readability. For effective visualization, raw attribution scores must be transformed into a standardized range. This normalized attribution intensity, which will later be mapped to visual properties such as color opacity (see Section 4.4), is defined as:

$$\text{intensity}(\mathbf{A}_{i, j}) = \frac{\mathbf{A}_{i, j} - \min(\mathbf{A})}{\max(\mathbf{A}) - \min(\mathbf{A})} \in [0, 1] \quad (7)$$

Beyond static visualization, the system must support dynamic exploration through context window adjustment, method toggling, token hovering interactions, model comparison, and linguistic category filtering. The context window size parameter  $c$  is constrained as  $1 \leq c \leq \lfloor \min(n, t)/2 \rfloor$ .

For multi-turn conversations, the follow-up input for turn  $k$  is constructed as:

$$X_k = (X_1, Y_1, X_2, Y_2, \dots, X_{k-1}, Y_{k-1}, X_k) \quad (8)$$

where each  $X_i$  and  $Y_i$  represent the user input and model output at turn  $i$ .

## 4 System Architecture and Visualization Methodology

This section introduces the conceptual framework for token attribution visualization, detailing its proposed architecture and key visualization methodologies. The proposed system is designed to address the requirements outlined in section 3.

## 4.1 System Overview and Backend Pipeline

The envisioned system provides an interactive visualization platform for exploring attribution information in sequence generation models, employing a job-based API architecture with Inseq as the backend attribution engine and featuring a reactive frontend for intuitive interaction with attribution data.

Using the formal definitions from section 3.1, the system operates on the fundamental sets  $\mathcal{M}$  (available models),  $\mathcal{F}$  (attribution methods),  $\mathcal{I}$  (possible input texts), and  $\mathcal{O}$  (method options). The computational unit is a job  $J = (\mathcal{M}', \mathcal{F}', i, o)$  as defined in Section 3.2, processing all combinations in the Cartesian product  $\mathcal{M}' \times \mathcal{F}'$ .

For each tuple  $(m, f) \in \mathcal{M}' \times \mathcal{F}'$ , the backend employs Inseq [43] to:

1. Load the model: `attribution_model = inseq.load_model(m, f)`
2. Compute attributions: `attributions = model.attribute(i, **o)`
3. Process and store results in the job’s result collection

The job-based API exposes four key endpoints: job creation (for initiating new attribution jobs), job status (for monitoring progress), job results (for retrieving completed analysis), and WebSocket notification (for real-time updates). The complete API specification is provided in Appendix A.1.

## 4.2 Result Data Structure

The result structure is designed to balance completeness with efficient transmission and rendering. Each job result includes job metadata, input text and tokenization, options, and individual results for each model-method combination. For each tuple  $(m, f)$ , the result contains model and method identifiers, generated output, raw attribution matrix  $\mathbf{A}^{m,f,X,Y} \in \mathbb{R}^{n \times t}$  as defined in section 3.1, pre-computed token importance aggregations  $I(x_i)$  and  $I(y_j)$  from section 3.3, and step scores for metrics such as probability and entropy.

The pre-computed aggregations are included to enable immediate visualization. The complete attribution matrix is attached for detailed exploration to be calculated client-side. The detailed schema is provided in Appendix A.2.

## 4.3 Frontend Interface and Visualization Techniques

The frontend comprises three primary views forming a coherent user flow:

The **Job Creation View** enables queuing of new attribution jobs through model selection, method selection, input text entry, and options configuration. Upon submission, the system creates a job with the specified parameters.

The **Job Overview** provides a centralized dashboard for managing running and completed jobs, with real-time status updates via WebSocket connection. When a job completes, it moves from Running to Completed status with a visual transition effect.

The **Attribution Dashboard** serves as the core interactive component for exploring attribution relationships. It features header controls (model selector, method toggles, context window slider, etc.), the main visualization area, and a follow-up job creation interface. This dashboard implements the visualization requirements specified in section 3.4. How the visualisation implements interactivity (which is a major benefit of the proposed system) is content of the following section 4.4.

## 4.4 Interactive Attribution Visualization

The core visualization interactively reveals attribution relationships between input and output tokens: For each attribution score  $\mathbf{A}_{i,j}$ , the visualization applies the normalization function defined in section 3.4. This normalized value is mapped to visual properties such as background color opacity.

The color encoding scheme assigns each attribution method a distinct color from a perceptually distinct palette, with appropriate blending techniques for combined visualization. A clear visual hierarchy maintains



text readability while conveying attribution information through background color, border style, and typography.

Four key interactive elements drive the exploration experience:

**Token Hovering:** When a user hovers over a token, the system highlights the hovered token and computes attribution relationships to all other tokens using the appropriate aggregation function from Section 3.3. This reveals specific token relationships dynamically.

**Context Window Control:** A slider adjusts the context window size parameter  $c$  within the constraints defined in Section 3.4. As this value changes, the system immediately updates the visualization using the corresponding aggregation functions.

**Method Selection:** Toggle buttons enable independent toggling of attribution methods, with visual feedback indicating active methods. When multiple methods are active, each is assigned a distinct color (e.g., red and blue), with overlapping attributions naturally creating blended visual effects (e.g., purple) in areas where both methods indicate importance.

**Model Comparison:** A split view toggle enables side-by-side comparison of different models, maintaining synchronized interactions between views for direct comparison of attribution patterns.

The default visualization presents each token colored by its overall importance ( $I(x_i)$  for input tokens and  $I(y_j)$  for output tokens) as defined in section 3.3, providing an immediate visual summary of influential tokens.

## 4.5 Advanced Visualization Capabilities

Beyond the core visualization, the system offers several advanced capabilities:

The before-mentioned toggle activates the **Model Comparison View**, which enables direct comparison of attribution patterns between different models. This split view synchronizes interactions across models and can optionally highlight differences through a differential attribution matrix  $\mathbf{A}_{i,j}^{\text{diff}} = \mathbf{A}_{i,j}^{m_1} - \mathbf{A}_{i,j}^{m_2}$ . Even when models generate different outputs for the same input, the system maintains appropriate alignment for comparison.

**Conversational Follow-up Analysis** traces attribution patterns through multi-turn conversations using the conversation construction formula from section 3.4. When viewing results from follow-up jobs, the system provides a conversation timeline, historical context highlighting showing how previous tokens influence current generation, and attribution flow analysis revealing evolving patterns across turns.

**Linguistic Category Filtering** enables analysis based on linguistic categories through integration with external NLP services. After enriching tokens with linguistic annotations, the system may support filtering by categories such as part-of-speech, named entities, and syntactic roles. Three visualization modes can be pictured: highlight mode (emphasizing selected categories), isolation mode (showing only selected categories), and comparison mode (displaying different categories with distinct visual encodings).

To ensure responsive interactions despite the computational intensity of attribution analysis, the system implements multiple optimization strategies: pre-computed aggregations for immediate visualization, on-demand client-side aggregation during interactions, throttled computations for larger matrices, and progressive loading for very large matrices.<sup>2</sup>

## 5 Implementation: TokenLense

With established requirements and architectural principles, we now bridge theory and practice by presenting TokenLense<sup>3</sup>, a proof-of-concept implementation of the proposed system.

### 5.1 System Implementation

TokenLense implements a three-component system that realizes key aspects of the conceptual framework:

---

<sup>2</sup>also see A.5

<sup>3</sup>TokenLense on GitHub

The **Extractor** component performs attribution analysis on language models using Inseq [43] to compute attribution matrices across multiple methods and models. It is implemented in Python and supports various attribution methods including attention weights, input  $\times$  gradient, integrated gradients, layer gradient  $\times$  activation, LIME, and saliency maps.


The **Analyzer** serves as a RESTful API layer implemented as a FastAPI service. While not implementing the complete job-based API architecture outlined in section 4.1, it provides critical endpoints for model discovery, method selection, and attribution data retrieval.

The **Visualizer** builds upon these endpoints to deliver interactive visualizations as detailed in section 4.3. The model selection interface allows users to browse available models and attribution methods, while the attribution dashboard provides the core visualization functionality.

## 5.2 Visualization Capabilities

The proof-of-concept implementation successfully realizes several key visualization capabilities defined above (section 4.4).

For the main visualisation of token importance (fig. 4), TokenLense supports two interactive modi operandi: The *first* one uses background highlighting, where hovering over a token causes the system to highlight related tokens with varying background color intensity based on attribution scores. The *second* one uses explicit connection lines between related tokens, with line thickness proportional to attribution strength.



(a) Entity identification visualization showing how the model processes a query about Warner Bros movies with a dog. Note the higher importance weights on key query terms.

(b) Named entity recognition example demonstrating how attribution patterns change during entity resolution, with varying importance weights assigned to contextual tokens.

**Figure 1** TokenLense token attribution visualization showing real-time importance analysis. The system uses background highlighting intensity to represent attribution scores, with orange highlighting for input tokens and blue highlighting for output tokens.

Figure 1 demonstrates the background highlighting approach, showing how token importance is visualized through color intensity. Additional examples of TokenLense visualizations are available in Appendix A.7.

TokenLense implements a comparison view (fig. 5) that enables side-by-side analysis of attribution patterns between different models or/and methods or/and aggregation methods.

## 5.3 Framework Alignment and Limitations

As a proof-of-concept, TokenLense implements substantial portions of the conceptual framework while necessarily omitting or simplifying others.

The core visualization requirements from section 3.4 are well-represented in the implementation, with TokenLense supporting interactive token relationship exploration, method selection, and model comparison through side-by-side visualization.

However, several aspects of the framework remain unimplemented in the current proof-of-concept. The job-based API architecture described in section 4.1 is simplified to a more traditional REST API without the comprehensive job management capabilities.

The implementation also demonstrates pragmatic departures from the theoretical framework in response to practical engineering concerns. For instance, rather than implementing real-time attribution computation, TokenLense uses pre-generated attribution data files to ensure responsive interactions.

## 5.4 Performance Considerations

TokenLens includes a performance analysis dashboard that extends the framework with quantitative evaluation capabilities. This component provides insights into the computational requirements of different attribution methods, addressing the efficiency considerations outlined in section 3.2.

## 5.5 Implementation Insights

The development and implementation of TokenLens yielded valuable insights while demonstrating the practical feasibility of key aspects of our attribution visualization framework. The successful implementation of interactive visualization techniques confirms the feasibility of real-time attribution exploration as proposed.

The implementation process revealed practical challenges that inform future development directions. The integration of diverse attribution methods highlights the need for standardized normalization approaches to enable meaningful comparison across methods with different numerical characteristics (see section 6).

As a proof-of-concept, TokenLens necessarily omits certain advanced features of the framework, particularly the comprehensive job-based architecture. However, it validates the fundamental approach of interactive attribution visualization while highlighting practical challenges that inform future refinements.

## 6 Future Work

### 6.1 Model and Method Support Extensions

The TokenLens implementation revealed significant challenges in applying attribution methods to a wide range of model architectures. While Inseq [43] provides a foundation for standardized attribution analysis, the practical implementation highlighted compatibility gaps within many model/method combinations, worsening with modern models (e.g. Gemma3 [46], Qwen [47] or Llama-4 [48] are not supported). Future work should prioritize extending attribution method implementations to support a wide architectural variations in models. This extension requires developing comprehensive testing frameworks that systematically validate attribution method compatibility across diverse model architectures.

Additionally, a robust model architecture detection system could automatically select compatible attribution methods based on the target model’s structure, avoiding failed attribution attempts that currently plague analysis workflows. This enhancement would address the **primary technical limitation** encountered in TokenLens — the inability to reliably analyze (newer) models with existing attribution methods—while maintaining the framework’s goal of **unified analysis** across diverse models.

### 6.2 Job-Based Architecture

The separation between computation and visualization components in TokenLens created workflow discontinuities that may significantly impact user experience. Implementing the job-based architecture outlined in Section 4.1 would transform the current multi-stage process into a seamless workflow, enabling asynchronous processing with real-time status updates via WebSockets while users continue exploration. Unified interfaces for job creation, monitoring, and result exploration would eliminate the current context-switching between separate components.

Efficient resource management for computationally intensive attribution methods represents a critical advancement, particularly for methods like Integrated Gradients that demand substantial processing power. Result caching and incremental visualization would further reduce latency, allowing immediate visualization of available results while computation continues in the background. These architectural enhancements directly address the practical engineering challenges identified during TokenLens development, creating a more integrated system that better supports exploratory analysis.

### 6.3 Conversational Analysis

Another limitation of the current implementation is its focus on single-turn analysis rather than multi-turn conversations that characterize modern LLM interactions. Building on TokenLense’s successful visualization of token relationships, future work should implement the conversational follow-up analysis capabilities described in Section 4.5. This extension would enable tracking attribution patterns throughout conversation history, revealing how previous exchanges influence current generation.

The visualization system could be enhanced to show how attribution patterns evolve across turns, identifying persistent patterns and emerging relationships. Supporting conversational context windows that span multiple exchanges would provide a more complete picture of model behavior in interactive settings. These capabilities are particularly important as LLMs usually operate in interactive contexts where understanding attribution across conversational history is essential for comprehensive interpretability, especially in applications like virtual assistants and customer service systems.

### 6.4 Advanced Analysis and Visualization

While TokenLense demonstrates the viability of interactive attribution visualization, future work should implement more sophisticated analysis capabilities. *Automated pattern recognition* algorithms could identify recurring attribution structures across inputs and models, building on the comparative visualization capabilities demonstrated in TokenLense. These algorithms would help users discover consistent patterns that might indicate model biases or specialized capabilities without requiring manual inspection of each input-output pair.

*Quantitative attribution metrics* would provide statistical measures for evaluating attribution quality and patterns, extending the performance analysis dashboard implemented in TokenLense to include quality metrics. Such metrics might quantify attribution concentration, stability across similar inputs, or alignment with linguistic structures, enabling more objective comparison between models or attribution methods.

Integration with *counterfactual analysis* would allow exploration of how input modifications affect attribution patterns, leveraging TokenLense’s model comparison capabilities to visualize before/after states. This would help identify which tokens are most sensitive to changes and how attribution patterns shift in response to different inputs. The performance analysis component of TokenLense provides a foundation for these advanced capabilities by demonstrating the feasibility of comparative analysis across methods and models.

TokenLense’s implementation of multiple visualization modes demonstrates the potential for diverse visual encodings. Building on this foundation, augmented reality (AR) could address the information density challenges encountered in two-dimensional interfaces. Three-dimensional spatial arrangement of tokens based on attribution relationships would enable more intuitive exploration of complex attribution patterns, allowing users to physically navigate through token spaces.

Gestural interactions for manipulating attribution parameters and navigating token relationships could make exploration more intuitive, particularly for users without technical expertise in model internals. Multi-dimensional visualization of attribution patterns not possible in traditional interfaces could reveal complex relationships obscured by two-dimensional projections. These approaches would leverage spatial cognition to create new ways of interacting with complex attribution data, potentially making attribution analysis more accessible to diverse stakeholders including domain experts without machine learning backgrounds.

### 6.5 Model Development Integration

The model comparison capabilities implemented in TokenLense revealed distinct attribution patterns across different models, validating the potential value of attribution visualization for model development. Future work should focus on tighter integration between attribution visualization and model development workflows. **Attribution-guided fine-tuning** represents a particularly promising direction, using identified attribution patterns to generate targeted training examples and monitor improvements during fine-tuning. This approach could help address specific weaknesses or biases revealed through attribution analysis.

**Prompt engineering support** would analyze how different prompt formulations affect attribution patterns, suggesting optimizations based on desired characteristics. This capability would be particularly valuable

as prompt engineering becomes increasingly important for extracting optimal performance from foundation models. Extending TokenLense’s comparison view to create comprehensive **model comparison frameworks** would enable systematic evaluation of attribution differences between model versions or architectures. These capabilities would transform attribution visualization from an analysis tool into an integral component of the model development process, addressing the broader goal of improving model transparency and reliability.

## 6.6 Implementation Roadmap

Based on the TokenLense experience, we propose a prioritized implementation roadmap that addresses the most critical limitations while progressively building toward more advanced capabilities. The first priority is extending Inseq’s model support to cover more recent architectures, as this represents the most immediate barrier to applying attribution analysis to state-of-the-art models. Following this, implementing the job-based architecture would address workflow discontinuity issues that currently hinder user experience.

With these foundations in place, adding conversation analysis support for multi-turn interactions would extend the system’s applicability to interactive contexts. Subsequently, developing quantitative attribution metrics and pattern recognition would enhance analytical capabilities, followed by integration with model development workflows to leverage attribution insights for model improvement. Finally, exploration of advanced visualization techniques, including augmented reality, would address information density challenges for complex models. This phased approach creates a clear development path from the current proof-of-concept toward the comprehensive framework envisioned in this paper.

## 7 Conclusion

This paper has proposed a comprehensive framework for visualizing and analyzing token attribution in large language models, combining formal mathematical foundations with interactive visualization techniques. The TokenLense implementation has demonstrated the practical feasibility of key components of this framework while revealing important insights for future development.

### 7.1 Framework and Implementation

Our work addresses critical gaps in existing interpretability tools: the inability to compare different attribution methods within a unified interface, the lack of interactive exploration capabilities, and the computational scalability challenges that current methods face. By formalizing a job-based processing architecture and interactive visualization interface, we provide a foundation for more comprehensive and nuanced analysis of token importance in modern LLMs.

The TokenLense implementation validates core aspects of this framework through several key accomplishments. The successful interactive visualization of token relationships through multiple visual encodings demonstrates that complex attribution information can be presented in an accessible format. The comparative analysis of attribution patterns across models and methods validates the framework’s emphasis on unified analysis capabilities, revealing distinct patterns that would be difficult to identify through separate tools. TokenLense also demonstrates practical visualization of complex attribution matrices, converting abstract numerical relationships into interpretable visual patterns. Additionally, the performance analysis dashboard provides valuable insights into computational requirements across attribution methods, informing future optimization efforts.

Implementation challenges revealed important areas for further development, like compatibility issues with newer model architectures and workflow discontinuities between separated components. The difficulty in applying existing attribution methods to given model architectures highlighted the need for more robust model compatibility, while the workflow discontinuities resulting from separated computation and visualization components underscored the importance of the proposed job-based architecture. The limited support for conversational analysis pointed to a critical gap in addressing modern interactive LLM applications, directly informing the future directions outlined in Section 6.

## 7.2 Implementation Priorities

Building on the TokenLense experience, future implementation efforts should proceed with clear priorities. Enhancing model compatibility represents the most urgent need, requiring extensions to attribution method support for newer architectures to enable analysis of state-of-the-art models. Following this foundation, implementing the complete job-based API architecture would transform the current multi-stage process into a seamless workflow with real-time feedback and efficient resource management.

With these core components in place, development should focus on conversational analysis capabilities to support multi-turn interactions, addressing a critical limitation in the current implementation. Finally, integration with model development workflows would enable attribution-guided improvement, transforming attribution visualization from an analysis tool into an integral component of the model development process. These priorities address the most significant limitations identified during implementation while maintaining the core vision of making complex model behavior more transparent and interpretable.

## 7.3 Broader Impact

The potential impact of this work extends beyond technical contributions to ML interpretability. In healthcare, the framework could help clinicians understand how medical LLMs weigh different symptoms in forming diagnostic suggestions, potentially identifying undesirable biases or unexpected prioritization patterns. In legal applications, it could reveal how models prioritize different precedents or statutes when generating analysis, providing critical transparency for judicial review and accountability. These domain-specific insights are essential for responsible deployment in high-stakes environments where understanding model reasoning is as important as the final output.

This research aligns with broader initiatives in AI transparency and governance, including emerging regulatory requirements that may soon mandate explainability for AI systems deployed in critical domains. The optimization strategies and job-based architecture are designed specifically to address computational feasibility concerns through efficient resource management and progressive loading techniques, making comprehensive attribution analysis practical even for resource-constrained environments.

## 7.4 Conclusion

The TokenLense implementation demonstrates that effective visualization of token attribution is both technically feasible and practically valuable for understanding LLM behavior. While challenges remain in scaling to billion-parameter models and supporting the full range of contemporary architectures, this work establishes a clear path toward more transparent and interpretable language models. The framework’s mathematical foundations provide theoretical rigor, while the interactive visualization approach ensures accessibility to diverse stakeholders.

By providing a rigorous yet accessible approach to understanding token importance, this framework represents a step toward LLMs that can be meaningfully audited, improved, and aligned with human values. The true value of this work will ultimately be measured by its ability to empower diverse stakeholders—from model developers to domain experts to regulators—with actionable insights into how these increasingly influential systems process and prioritize information. As LLMs continue to transform society through their integration into critical systems, such transparency tools will become essential infrastructure for responsible AI development and deployment.

## A Appendix

### A.1 Framework API Specification

This appendix provides the specification for the system's REST API, enabling programmatic interaction with the attribution visualization framework.

**Job Creation Endpoint** Creates a new attribution job with specified models, methods, input, and options.

```
POST /api/jobs
{
  "models": ["model_1", "model_2", ...],
  "methods": ["method_1", "method_2", ...],
  "input": "input text",
  "options": { ... }
}
```

**Job Status Endpoint** Provides the current status of a job.

```
GET /api/jobs/{job_id}/status
```

**Job Results Endpoint** Returns complete results when a job is completed.

```
GET /api/jobs/{job_id}/results
```

**WebSocket Notification** Sends real-time updates on job status changes.

```
WS /api/ws/jobs/{job_id}
```

### A.2 Framework Result Data Structure Schema

The complete job result structure follows this schema:

```
{
  "job_id": "unique_job_id",
  "created_at": "ISO datetime",
  "completed_at": "ISO datetime",
  "input": {
    "text": "original input text",
    "tokens": ["x_1", "x_2", ..., "x_n"]
  },
  "options": { ... },
  "results": [
    {
      "model": "model_name",
      "method": "method_name",
      "output": {
        "text": "generated output text",
        "tokens": ["y_1", "y_2", ..., "y_t"]
      },
      "attribution": {
        "raw": [ ... ], // Attribution matrix
        "aggregated": {
          "token_importance": { ... }
        }
      }
    }
  ]
}
```

```

    }
  },
  "step_scores": { ... }
},
...
]
}

```

### A.3 Advanced Aggregation Functions

In addition to the core aggregation functions described in Section 3.3, the system implements several advanced aggregation methods:

**Dimensional Aggregation** For multi-dimensional attribution methods:

$$\mathbf{A}_{i,j} = \text{aggregate}(\mathbf{A}_{i,j,1}, \mathbf{A}_{i,j,2}, \dots, \mathbf{A}_{i,j,d}) \quad (9)$$

Where aggregation can be L2 Norm:  $\sqrt{\sum_{k=1}^d \mathbf{A}_{i,j,k}^2}$ , Sum:  $\sum_{k=1}^d \mathbf{A}_{i,j,k}$ , or Mean:  $\frac{1}{d} \sum_{k=1}^d \mathbf{A}_{i,j,k}$ .

**Multi-method Aggregation** When multiple methods are active:

$$\mathbf{A}_{i,j}^{\text{combined}} = \frac{1}{|F'|} \sum_{f \in F'} \mathbf{A}_{i,j}^f \quad (10)$$

**Temporal Aggregation** For analyzing attribution across conversation turns:

$$\mathbf{A}_{i,j,k}^{\text{temporal}} = \sum_{t=1}^k \alpha_t \cdot \mathbf{A}_{i,j}^t \quad (11)$$

Where  $\alpha_t$  is a weighting factor for turn  $t$  and  $\mathbf{A}_{i,j}^t$  is the attribution score at turn  $t$ .

### A.4 Frontend Views

**Job Creation View** Includes model selection (multi-select with metadata), method selection (with compatibility indicators), input text area (with token visualization), options configuration (with method-specific options), template selection, and submit button.

**Job Overview** Provides running jobs table (with progress indicators), completed jobs table (with status and view buttons), filter controls, sort controls, search field, and batch operations.

**Attribution Dashboard** Features header controls, visualization area (with hover, click, and drag interactions), attribution details panel, method comparison panel, token statistics, export controls, and follow-up interface.

### A.5 Performance Optimizations

**Parallel Processing** Implemented through task queue architecture with priority-based scheduling, worker pool management with dynamic scaling, resource allocation algorithms, and fine-grained task decomposition.

**Caching Strategy** Multi-level caching includes model cache (loaded models with LRU eviction), tokenization cache (indexed by text and model), attribution cache (with content-based hashing), and aggregation cache (pre-computed common aggregations).



**Progressive Loading** Divides large attribution matrices into tiles based on viewport, implements priority loading based on user interactions, uses background loading during idle periods, and maintains a local cache of loaded tiles.

## A.6 Implementation Technologies


**Backend Technologies** The system builds upon Inseq for attribution computation, wrapped in a job-based API structure. The backend uses a modular architecture with a job management service, attribution computation engine, API and WebSocket service, and results database. Performance optimizations include parallel processing, caching, resource management, and compressed data transmission.

**Frontend Technologies** The React/TypeScript application features a layered component architecture with core visualization components, specialized view components, and layout components. State management uses a structured approach with global application state, component-local state, and derived state. Visualization optimizations include virtualized rendering, throttled event handlers, web workers for heavy computation, memoization, and efficient DOM updates.

**Inseq Integration** The system utilizes Inseq's core functions for model loading and attribution computation while extending its visualization capabilities with interactive features. An optimization layer pre-computes common aggregations, optimizes memory usage, implements caching strategies, and provides progressive loading for large datasets.

## A.7 Visualization Examples

This section presents additional visualization examples from the TokenLense implementation, showcasing how the framework renders token attribution patterns across different query types and reasoning tasks.



**Figure 2** Background highlighting intensity corresponds to token importance in both input and output sequences, revealing how the model allocates attention during different reasoning tasks. These visualizations demonstrate the framework's ability to expose underlying patterns in model inference across various query types.


The examples above illustrate how TokenLense visualizes token importance weights during different types of language model inference:

1. Example 1 demonstrates the model’s entity search behavior, with higher importance weights on key search terms ("what", "movie", "has", "dog") as it formulates a strategy to identify movies produced by Warner Bros featuring a dog.
2. Example 2 shows geographical entity resolution, highlighting the model’s reasoning process as it considers different regions (Middle East versus Far East) for the ambiguous input.
3. Example 3 visualizes attribution patterns during named entity recognition for Detroit musicians, showing how the model structures its search parameters and maintains context.
4. Example 4 displays attribution during factual recall, with emphasis on descriptive tokens that characterize the entity being identified.

These visualizations complement the performance-focused screenshots in Appendix A.8 by focusing on the token-level attribution patterns that emerge during various reasoning tasks.

## A.8 Implementation Screenshots

This section presents screenshots from the TokenLense proof-of-concept implementation, showcasing various visualization features and interfaces.




**Figure 3** Model attribution method selection interface for the DeepSeek-R1-Distill-Qwen-7B model. The interface provides descriptions for different attribution methods including layer gradient  $\times$  activation, saliency, integrated gradients, input  $\times$  gradient, and attention.


## A.9 Framework and Implementation Comparison

Table 2 provides a detailed comparison between the conceptual framework proposed in Section 4 and the TokenLense proof-of-concept implementation described in Section 5.

This comparison highlights that TokenLense successfully implements many core visualization capabilities from the conceptual framework while simplifying certain architectural elements, particularly the job-based API structure. These differences reflect the practical constraints of proof-of-concept development **rather than limitations in the conceptual framework** itself.




**Figure 4** Interactive token highlighting interface showcasing token relationships with configurable aggregation methods (sum, mean, L2 norm, absolute sum, max) and visualization modes (background highlighting or connection lines).




**Figure 5** Comparative visualization of input-output token relationships between two different Llama models (Llama-3-2-1B and Llama-3-1-8B). The display shows how the same input prompt produces different token progressions in the output.

## Tokens Processed Per Second


Heatmap visualization of tokens processed per second for each model and attribution method



**Figure 6** Performance heatmap visualization displaying tokens processed per second for each model and attribution method combination. The logarithmic scale highlights performance differences across methods.



**Figure 7** Method performance comparison showing average processing time across different attribution methods and models. The visualization allows grouping by model and sorting by performance metrics.



**Figure 8** Token importance visualization showing tokens ranked by their importance across analyzed files. The visualization presents tokens with their frequency metrics and relative importance scores.

**Table 2** Comparison between Conceptual Framework and TokenLense Implementation

| Component                  | Conceptual Framework                                                                     | TokenLense Implementation                                                |
|----------------------------|------------------------------------------------------------------------------------------|--------------------------------------------------------------------------|
| <b>Architecture</b>        |                                                                                          |                                                                          |
| System Structure           | Job-based API architecture with reactive frontend                                        | Three-component system with simplified workflow                          |
| Processing Model           | Dynamic job creation with Cartesian product $\mathcal{M}' \times \mathcal{F}'$ execution | Batch processing of attribution methods without unified job abstraction  |
| Data Flow                  | Automated pipeline with WebSocket notifications                                          | Manual file transfer between components                                  |
| <b>Attribution Methods</b> |                                                                                          |                                                                          |
| Method Support             | All methods in Inseq with unified processing                                             | All methods in Inseq with unified processing, requires manual execution  |
| <b>API Implementation</b>  |                                                                                          |                                                                          |
| Endpoints                  | Job creation, status, results, and WebSocket notification                                | REST API for model discovery, method selection, and data retrieval       |
| Updates                    | Real-time WebSocket-based notification system                                            | Request-based with no real-time updates                                  |
| <b>Visualization</b>       |                                                                                          |                                                                          |
| Token Relationships        | Multiple techniques with advanced interaction models                                     | Background highlighting and connection line visualizations               |
| Interactive Elements       | Token hovering, context window, method selection, model comparison                       | Token hovering/selection, method toggling, multi method/model comparison |
| Advanced Features          | Differential attribution matrix, conversational follow-up, linguistic filtering          | Side-by-side comparison without differential analysis                    |
| <b>Data Processing</b>     |                                                                                          |                                                                          |
| Result Structure           | Complete structure with all metadata and matrices                                        | Simplified structure focusing on attribution data                        |
| Aggregation Functions      | Dimensional, multi-method, and temporal aggregations                                     | Inseq supported aggregations: sum, mean, L2 norm, abs sum, max           |
| <b>Performance</b>         |                                                                                          |                                                                          |
| Backend                    | Parallel processing, multi-level caching, resource allocation                            | No caching                                                               |
| Frontend                   | Virtualized rendering, progressive loading, background computation                       | No caching                                                               |

## References

- [1] L. De Angelis, F. Baglivo, G. Arzilli, G. P. Privitera, P. Ferragina, A. E. Tozzi, and C. Rizzo, “Chatgpt and the rise of large language models: the new ai-driven infodemic threat in public health,” *Frontiers in public health*, vol. 11, p. 1166120, 2023.
- [2] K. Carolan, L. Fennelly, and A. F. Smeaton, “A review of multi-modal large language and vision models,” *ArXiv*, vol. abs/2404.01322, 2024.
- [3] M. Ferdaus, M. Abdelguerfi, E. Ioup, K. N. Niles, K. Pathak, and S. Sloan, “Towards trustworthy ai: A review of ethical and robust large language models,” *ArXiv*, vol. abs/2407.13934, 2024.
- [4] M. A. Babu, K. M. Yusuf, L. N. Eni, S. M. S. Jaman, and M. R. Sharmin, “Chatgpt and generation ‘z’: A study on the usage rates of chatgpt,” *Social Sciences & Humanities Open*, vol. 10, p. 101163, 2024.
- [5] A. Humlum and E. Vestergaard, “The adoption of chatgpt,” IZA Discussion Papers 16992, Bonn, 2024.
- [6] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. u. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems* (I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, eds.), vol. 30, Curran Associates, Inc., 2017.
- [7] A. Radford, K. Narasimhan, T. Salimans, I. Sutskever, *et al.*, “Improving language understanding by generative pre-training,” 2018.
- [8] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)* (J. Burstein, C. Doran, and T. Solorio, eds.), (Minneapolis, Minnesota), pp. 4171–4186, Association for Computational Linguistics, June 2019.
- [9] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei, “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems* (H. Larochelle, M. Ranzato, R. Hadsell, M. Balcan, and H. Lin, eds.), vol. 33, pp. 1877–1901, Curran Associates, Inc., 2020.
- [10] OpenAI, J. Achiam, S. Adler, S. Agarwal, L. Ahmad, I. Akkaya, F. L. Aleman, D. Almeida, J. Altenschmidt, S. Altman, S. Anadkat, R. Avila, I. Babuschkin, S. Balaji, V. Balcom, P. Baltescu, H. Bao, M. Bavarian, J. Belgum, I. Bello, J. Berdine, G. Bernadett-Shapiro, C. Berner, L. Bogdonoff, O. Boiko, M. Boyd, A.-L. Brakman, G. Brockman, T. Brooks, M. Brundage, K. Button, T. Cai, R. Campbell, A. Cann, B. Carey, C. Carlson, R. Carmichael, B. Chan, C. Chang, F. Chantzis, D. Chen, S. Chen, R. Chen, J. Chen, M. Chen, B. Chess, C. Cho, C. Chu, H. W. Chung, D. Cummings, J. Currier, Y. Dai, C. Decareaux, T. Degry, N. Deutsch, D. Deville, A. Dhar, D. Dohan, S. Dowling, S. Dunning, A. Ecoffet, A. Eleti, T. Eloundou, D. Farhi, L. Fedus, N. Felix, S. P. Fishman, J. Forte, I. Fulford, L. Gao, E. Georges, C. Gibson, V. Goel, T. Gogineni, G. Goh, R. Gontijo-Lopes, J. Gordon, M. Grafstein, S. Gray, R. Greene, J. Gross, S. S. Gu, Y. Guo, C. Hallacy, J. Han, J. Harris, Y. He, M. Heaton, J. Heidecke, C. Hesse, A. Hickey, W. Hickey, P. Hoeschele, B. Houghton, K. Hsu, S. Hu, X. Hu, J. Huizinga, S. Jain, S. Jain, J. Jang, A. Jiang, R. Jiang, H. Jin, D. Jin, S. Jomoto, B. Jonn, H. Jun, T. Kaftan, Łukasz Kaiser, A. Kamali, I. Kanitscheider, N. S. Keskar, T. Khan, L. Kilpatrick, J. W. Kim, C. Kim, Y. Kim, J. H. Kirchner, J. Kiros, M. Knight, D. Kokotajlo, Łukasz Kondraciuk, A. Kondrich, A. Konstantinidis, K. Kosic, G. Krueger, V. Kuo, M. Lampe, I. Lan, T. Lee, J. Leike, J. Leung, D. Levy, C. M. Li, R. Lim, M. Lin, S. Lin, M. Litwin, T. Lopez, R. Lowe, P. Lue, A. Makanju, K. Malfacini, S. Manning, T. Markov, Y. Markovski, B. Martin, K. Mayer, A. Mayne, B. McGrew, S. M. McKinney, C. McLeavey, P. McMillan, J. McNeil, D. Medina, A. Mehta, J. Menick,

L. Metz, A. Mishchenko, P. Mishkin, V. Monaco, E. Morikawa, D. Mossing, T. Mu, M. Murati, O. Murk, D. Mély, A. Nair, R. Nakano, R. Nayak, A. Neelakantan, R. Ngo, H. Noh, L. Ouyang, C. O’Keefe, J. Pachocki, A. Paino, J. Palermo, A. Pantuliano, G. Parascandolo, J. Parish, E. Parparita, A. Passos, M. Pavlov, A. Peng, A. Perelman, F. de Avila Belbute Peres, M. Petrov, H. P. de Oliveira Pinto, Michael, Pokorny, M. Pokrass, V. H. Pong, T. Powell, A. Power, B. Power, E. Proehl, R. Puri, A. Radford, J. Rae, A. Ramesh, C. Raymond, F. Real, K. Rimbach, C. Ross, B. Rotsted, H. Roussez, N. Ryder, M. Saltarelli, T. Sanders, S. Santurkar, G. Sastry, H. Schmidt, D. Schnurr, J. Schulman, D. Selsam, K. Sheppard, T. Sherbakov, J. Shieh, S. Shoker, P. Shyam, S. Sidor, E. Sigler, M. Simens, J. Sitkin, K. Slama, I. Sohl, B. Sokolowsky, Y. Song, N. Staudacher, F. P. Such, N. Summers, I. Sutskever, J. Tang, N. Tezak, M. B. Thompson, P. Tillet, A. Tootoonchian, E. Tseng, P. Tuggle, N. Turley, J. Tworek, J. F. C. Uribe, A. Vallone, A. Vijayvergiya, C. Voss, C. Wainwright, J. J. Wang, A. Wang, B. Wang, J. Ward, J. Wei, C. Weinmann, A. Welihinda, P. Welinder, J. Weng, L. Weng, M. Wiethoff, D. Willner, C. Winter, S. Wolrich, H. Wong, L. Workman, S. Wu, J. Wu, M. Wu, K. Xiao, T. Xu, S. Yoo, K. Yu, Q. Yuan, W. Zaremba, R. Zellers, C. Zhang, M. Zhang, S. Zhao, T. Zheng, J. Zhuang, W. Zhuk, and B. Zoph, “Gpt-4 technical report,” 2024.

- [11] A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Vaughan, A. Yang, A. Fan, A. Goyal, A. Hartshorn, A. Yang, A. Mitra, A. Sravankumar, A. Korenev, A. Hinsvark, A. Rao, A. Zhang, A. Rodriguez, A. Gregerson, A. Spataru, B. Roziere, B. Biron, B. Tang, B. Chern, C. Caucheteux, C. Nayak, C. Bi, C. Marra, C. McConnell, C. Keller, C. Touret, C. Wu, C. Wong, C. C. Ferrer, C. Nikolaidis, D. Allonsius, D. Song, D. Pintz, D. Livshits, D. Wyatt, D. Esiobu, D. Choudhary, D. Mahajan, D. Garcia-Olano, D. Perino, D. Hupkes, E. Lakomkin, E. AlBadawy, E. Lobanova, E. Dinan, E. M. Smith, F. Radenovic, F. Guzmán, F. Zhang, G. Synnaeve, G. Lee, G. L. Anderson, G. Thattai, G. Nail, G. Mialon, G. Pang, G. Cucurell, H. Nguyen, H. Korevaar, H. Xu, H. Touvron, I. Zarov, I. A. Ibarra, I. Kloumann, I. Misra, I. Evtimov, J. Zhang, J. Copet, J. Lee, J. Geffert, J. Vranes, J. Park, J. Mahadeokar, J. Shah, J. van der Linde, J. Billock, J. Hong, J. Lee, J. Fu, J. Chi, J. Huang, J. Liu, J. Wang, J. Yu, J. Bitton, J. Spisak, J. Park, J. Rocca, J. Johnstun, J. Saxe, J. Jia, K. V. Alwala, K. Prasad, K. Upasani, K. Plawiak, K. Li, K. Heafield, K. Stone, K. El-Arini, K. Iyer, K. Malik, K. Chiu, K. Bhalla, K. Lakhota, L. Rantala-Yearly, L. van der Maaten, L. Chen, L. Tan, L. Jenkins, L. Martin, L. Madaan, L. Malo, L. Blecher, L. Landzaat, L. de Oliveira, M. Muzzi, M. Pasupuleti, M. Singh, M. Paluri, M. Kardas, M. Tsimpoukelli, M. Oldham, M. Rita, M. Pavlova, M. Kambadur, M. Lewis, M. Si, M. K. Singh, M. Hassan, N. Goyal, N. Torabi, N. Bashlykov, N. Bogoychev, N. Chatterji, N. Zhang, O. Duchenne, O. Çelebi, P. Alrassy, P. Zhang, P. Li, P. Vasic, P. Weng, P. Bhargava, P. Dubal, P. Krishnan, P. S. Koura, P. Xu, Q. He, Q. Dong, R. Srinivasan, R. Ganapathy, R. Calderer, R. S. Cabral, R. Stojnic, R. Raileanu, R. Maheswari, R. Girdhar, R. Patel, R. Sauvestre, R. Polidoro, R. Sumbaly, R. Taylor, R. Silva, R. Hou, R. Wang, S. Hosseini, S. Chennabasappa, S. Singh, S. Bell, S. S. Kim, S. Edunov, S. Nie, S. Narang, S. Raparthy, S. Shen, S. Wan, S. Bhosale, S. Zhang, S. Vandenhende, S. Batra, S. Whitman, S. Sootla, S. Collot, S. Gururangan, S. Borodinsky, T. Herman, T. Fowler, T. Sheasha, T. Georgiou, T. Scialom, T. Speckbacher, T. Mihaylov, T. Xiao, U. Karn, V. Goswami, V. Gupta, V. Ramanathan, V. Kerkez, V. Gonguet, V. Do, V. Vogeti, V. Albiero, V. Petrovic, W. Chu, W. Xiong, W. Fu, W. Meers, X. Martinet, X. Wang, X. Wang, X. E. Tan, X. Xia, X. Xie, X. Jia, X. Wang, Y. Goldschlag, Y. Gaur, Y. Babaei, Y. Wen, Y. Song, Y. Zhang, Y. Li, Y. Mao, Z. D. Coudert, Z. Yan, Z. Chen, Z. Papakipos, A. Singh, A. Srivastava, A. Jain, A. Kelsey, A. Shajnfeld, A. Gangidi, A. Victoria, A. Goldstand, A. Menon, A. Sharma, A. Boesenberg, A. Baevski, A. Feinstein, A. Kallet, A. Sangani, A. Teo, A. Yunus, A. Lupu, A. Alvarado, A. Caples, A. Gu, A. Ho, A. Poulton, A. Ryan, A. Ramchandani, A. Dong, A. Franco, A. Goyal, A. Saraf, A. Chowdhury, A. Gabriel, A. Bharambe, A. Eisenman, A. Yazdan, B. James, B. Maurer, B. Leonhardi, B. Huang, B. Loyd, B. D. Paola, B. Paranjape, B. Liu, B. Wu, B. Ni, B. Hancock, B. Wasti, B. Spence, B. Stojkovic, B. Gamido, B. Montalvo, C. Parker, C. Burton, C. Mejia, C. Liu, C. Wang, C. Kim, C. Zhou, C. Hu, C.-H. Chu, C. Cai, C. Tindal, C. Feichtenhofer, C. Gao, D. Civin, D. Beaty, D. Kreymer, D. Li, D. Adkins, D. Xu, D. Testuggine, D. David, D. Parikh, D. Liskovich, D. Foss, D. Wang, D. Le, D. Holland, E. Dowling, E. Jamil, E. Montgomery, E. Presani, E. Hahn, E. Wood, E.-T. Le, E. Brinkman, E. Arcaute, E. Dunbar, E. Smothers, F. Sun, F. Kreuk, F. Tian, F. Kokkinos, F. Ozgenel, F. Caggioni, F. Kanayet, F. Seide, G. M. Florez, G. Schwarz, G. Badeer, G. Swee, G. Halpern, G. Herman, G. Sizov, Guangyi, Zhang,

G. Lakshminarayanan, H. Inan, H. Shojanazeri, H. Zou, H. Wang, H. Zha, H. Habeeb, H. Rudolph, H. Suk, H. Aspegren, H. Goldman, H. Zhan, I. Damlaj, I. Molybog, I. Tufanov, I. Leontiadis, I.-E. Veliche, I. Gat, J. Weissman, J. Geboski, J. Kohli, J. Lam, J. Asher, J.-B. Gaya, J. Marcus, J. Tang, J. Chan, J. Zhen, J. Reizenstein, J. Teboul, J. Zhong, J. Jin, J. Yang, J. Cummings, J. Carvill, J. Shepard, J. McPhie, J. Torres, J. Ginsburg, J. Wang, K. Wu, K. H. U, K. Saxena, K. Khandelwal, K. Zand, K. Matosich, K. Veeraraghavan, K. Michelena, K. Li, K. Jagadeesh, K. Huang, K. Chawla, K. Huang, L. Chen, L. Garg, L. A, L. Silva, L. Bell, L. Zhang, L. Guo, L. Yu, L. Moshkovich, L. Wehrstedt, M. Khabsa, M. Avalani, M. Bhatt, M. Mankus, M. Hasson, M. Lennie, M. Reso, M. Groshev, M. Naumov, M. Lathi, M. Keneally, M. Liu, M. L. Seltzer, M. Valko, M. Restrepo, M. Patel, M. Vyatskov, M. Samvelyan, M. Clark, M. Macey, M. Wang, M. J. Hermoso, M. Metanat, M. Rastegari, M. Bansal, N. Santhanam, N. Parks, N. White, N. Bawa, N. Singhal, N. Egebo, N. Usunier, N. Mehta, N. P. Laptev, N. Dong, N. Cheng, O. Chernoguz, O. Hart, O. Salpekar, O. Kalinli, P. Kent, P. Parekh, P. Saab, P. Balaji, P. Rittner, P. Bontrager, P. Roux, P. Dollar, P. Zvyagina, P. Ratanchandani, P. Yuvraj, Q. Liang, R. Alao, R. Rodriguez, R. Ayub, R. Murthy, R. Nayani, R. Mitra, R. Parthasarathy, R. Li, R. Hogan, R. Battey, R. Wang, R. Howes, R. Rinott, S. Mehta, S. Siby, S. J. Bondu, S. Datta, S. Chugh, S. Hunt, S. Dhillon, S. Sidorov, S. Pan, S. Mahajan, S. Verma, S. Yamamoto, S. Ramaswamy, S. Lindsay, S. Lindsay, S. Feng, S. Lin, S. C. Zha, S. Patil, S. Shankar, S. Zhang, S. Zhang, S. Wang, S. Agarwal, S. Sajuyigbe, S. Chintala, S. Max, S. Chen, S. Kehoe, S. Satterfield, S. Govindaprasad, S. Gupta, S. Deng, S. Cho, S. Virk, S. Subramanian, S. Choudhury, S. Goldman, T. Remez, T. Glaser, T. Best, T. Koehler, T. Robinson, T. Li, T. Zhang, T. Matthews, T. Chou, T. Shaked, V. Vontimitta, V. Ajayi, V. Montanez, V. Mohan, V. S. Kumar, V. Mangla, V. Ionescu, V. Poenaru, V. T. Mihailescu, V. Ivanov, W. Li, W. Wang, W. Jiang, W. Bouaziz, W. Constable, X. Tang, X. Wu, X. Wang, X. Wu, X. Gao, Y. Kleinman, Y. Chen, Y. Hu, Y. Jia, Y. Qi, Y. Li, Y. Zhang, Y. Zhang, Y. Adi, Y. Nam, Yu, Wang, Y. Zhao, Y. Hao, Y. Qian, Y. Li, Y. He, Z. Rait, Z. DeVito, Z. Rosnbrick, Z. Wen, Z. Yang, Z. Zhao, and Z. Ma, “The llama 3 herd of models,” 2024.

- [12] M. Truell, W. Yan, J. Palmer, M. Catasta, D. Wu, C. Falls, A. Li, Y. Bhavnani, C. Allen, S. Last, and et al., “Claude 3.7 sonnet.”
- [13] A. J. Thirunavukarasu, D. S. J. Ting, K. Elangovan, L. Gutierrez, T. F. Tan, and D. S. W. Ting, “Large language models in medicine,” *Nature Medicine*, vol. 29, pp. 1930–1940, Aug 2023.
- [14] S. Wu, O. Irsoy, S. Lu, V. Dabrovolski, M. Dredze, S. Gehrmann, P. Kambadur, D. Rosenberg, and G. Mann, “Bloomberggpt: A large language model for finance,” *arXiv preprint arXiv:2303.17564*, 2023.
- [15] A. Abd-alrazaq, R. AlSaad, D. Alhuwail, A. Ahmed, P. M. Healy, S. Latifi, S. Aziz, R. Damseh, S. Alabed Alrazak, and J. Sheikh, “Large language models in medical education: Opportunities, challenges, and future directions,” *JMIR Med Educ*, vol. 9, p. e48291, Jun 2023.
- [16] J. Cui, Z. Li, Y. Yan, B. Chen, and L. Yuan, “Chatlaw: Open-source legal large language model with integrated external knowledge bases,” *CoRR*, vol. abs/2306.16092, 2023.
- [17] R. Bommasani, D. A. Hudson, E. Adeli, R. Altman, S. Arora, S. von Arx, M. S. Bernstein, J. Bohg, A. Bosselut, E. Brunskill, E. Brynjolfsson, S. Buch, D. Card, R. Castellon, N. Chatterji, A. Chen, K. Creel, J. Q. Davis, D. Demszky, C. Donahue, M. Doumbouya, E. Durmus, S. Ermon, J. Etchemendy, K. Ethayarajh, L. Fei-Fei, C. Finn, T. Gale, L. Gillespie, K. Goel, N. Goodman, S. Grossman, N. Guha, T. Hashimoto, P. Henderson, J. Hewitt, D. E. Ho, J. Hong, K. Hsu, J. Huang, T. Icard, S. Jain, D. Jurafsky, P. Kalluri, S. Karamcheti, G. Keeling, F. Khani, O. Khattab, P. W. Koh, M. Krass, R. Krishna, R. Kuditipudi, A. Kumar, F. Ladhak, M. Lee, T. Lee, J. Leskovec, I. Levent, X. L. Li, X. Li, T. Ma, A. Malik, C. D. Manning, S. Mirchandani, E. Mitchell, Z. Munyikwa, S. Nair, A. Narayan, D. Narayanan, B. Newman, A. Nie, J. C. Niebles, H. Nilforoshan, J. Nyarko, G. Ogut, L. Orr, I. Papadimitriou, J. S. Park, C. Piech, E. Portelance, C. Potts, A. Raghunathan, R. Reich, H. Ren, F. Rong, Y. Roohani, C. Ruiz, J. Ryan, C. Ré, D. Sadigh, S. Sagawa, K. Santhanam, A. Shih, K. Srinivasan, A. Tamkin, R. Taori, A. W. Thomas, F. Tramèr, R. E. Wang, W. Wang, B. Wu, J. Wu, Y. Wu, S. M. Xie, M. Yasunaga, J. You, M. Zaharia,



- M. Zhang, T. Zhang, X. Zhang, Y. Zhang, L. Zheng, K. Zhou, and P. Liang, “On the opportunities and risks of foundation models,” 2022.
- [18] L. Weidinger, J. Uesato, M. Rauh, C. Griffin, P.-S. Huang, J. Mellor, A. Glaese, M. Cheng, B. Balle, A. Kasirzadeh, C. Biles, S. Brown, Z. Kenton, W. Hawkins, T. Stepleton, A. Birhane, L. A. Hendricks, L. Rimell, W. Isaac, J. Haas, S. Legassick, G. Irving, and I. Gabriel, “Taxonomy of risks posed by language models,” in *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’22, (New York, NY, USA), p. 214–229, Association for Computing Machinery, 2022.
- [19] S. R. Bowman, “Eight things to know about large language models,” 2023.
- [20] J. Kaplan, S. McCandlish, T. Henighan, T. B. Brown, B. Chess, R. Child, S. Gray, A. Radford, J. Wu, and D. Amodei, “Scaling laws for neural language models,” *arXiv preprint arXiv:2001.08361*, 2020.
- [21] Q. V. Liao and J. W. Vaughan, “Ai transparency in the age of llms: A human-centered research roadmap,” *arXiv preprint arXiv:2306.01941*, vol. 10, 2023.
- [22] E. Cambria, L. Malandri, F. Mercurio, N. Nobani, and A. Seveso, “Xai meets llms: A survey of the relation between explainable ai and large language models,” 2024.
- [23] H. Zhao, H. Chen, F. Yang, N. Liu, H. Deng, H. Cai, S. Wang, D. Yin, and M. Du, “Explainability for large language models: A survey,” *ACM Trans. Intell. Syst. Technol.*, vol. 15, Feb. 2024.
- [24] A. Bibal, R. Cardon, D. Alfter, R. Wilkens, X. Wang, T. François, and P. Watrin, “Is attention explanation? an introduction to the debate,” in *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (S. Muresan, P. Nakov, and A. Villavicencio, eds.), (Dublin, Ireland), pp. 3889–3900, Association for Computational Linguistics, May 2022.
- [25] N. Elhage, N. Nanda, C. Olsson, T. Henighan, N. Joseph, B. Mann, A. Askell, Y. Bai, A. Chen, T. Conerly, N. DasSarma, D. Drain, D. Ganguli, Z. Hatfield-Dodds, D. Hernandez, A. Jones, J. Kernion, L. Lovitt, K. Ndousse, D. Amodei, T. Brown, J. Clark, J. Kaplan, S. McCandlish, and C. Olah, “A mathematical framework for transformer circuits,” *Transformer Circuits*, Dec 2021. Last accessed March 11, 2025.
- [26] O. Antverg and Y. Belinkov, “On the pitfalls of analyzing individual neurons in language models,” 2022.
- [27] M. Turpin, J. Michael, E. Perez, and S. Bowman, “Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting,” in *Advances in Neural Information Processing Systems* (A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, eds.), vol. 36, pp. 74952–74965, Curran Associates, Inc., 2023.
- [28] A. Rogers, O. Kovaleva, and A. Rumshisky, “A primer in bertology: What we know about how bert works,” *Transactions of the Association for Computational Linguistics*, vol. 8, pp. 842–866, 01 2021.
- [29] M. Sundararajan, A. Taly, and Q. Yan, “Axiomatic attribution for deep networks,” in *Proceedings of the 34th International Conference on Machine Learning* (D. Precup and Y. W. Teh, eds.), vol. 70 of *Proceedings of Machine Learning Research*, pp. 3319–3328, PMLR, 06–11 Aug 2017.
- [30] M. T. Ribeiro, S. Singh, and C. Guestrin, “Model-agnostic interpretability of machine learning,” 2016.
- [31] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems* (I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, eds.), vol. 30, Curran Associates, Inc., 2017.
- [32] J. Vig, “A multiscale visualization of attention in the transformer model,” 2019.
- [33] S. Jain and B. C. Wallace, “Attention is not explanation,” 2019.
- [34] S. Wiegrefe and Y. Pinter, “Attention is not not explanation,” 2019.

- [35] I. Tenney, D. Das, and E. Pavlick, “Bert rediscovers the classical nlp pipeline,” 2019.
- [36] B. Kim, M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viegas, and R. sayres, “Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (TCAV),” in *Proceedings of the 35th International Conference on Machine Learning* (J. Dy and A. Krause, eds.), vol. 80 of *Proceedings of Machine Learning Research*, pp. 2668–2677, PMLR, 10–15 Jul 2018.
- [37] A. Bau, Y. Belinkov, H. Sajjad, N. Durrani, F. Dalvi, and J. Glass, “Identifying and controlling important neurons in neural machine translation,” 2018.
- [38] M. Weiser, “The computer for the 21 st century,” *Scientific american*, vol. 265, no. 3, pp. 94–105, 1991.
- [39] E. Dimara and C. Perin, “What is interaction for data visualization?,” *IEEE Transactions on Visualization and Computer Graphics*, vol. 26, no. 1, pp. 119–129, 2020.
- [40] R. Sennrich, B. Haddow, and A. Birch, “Neural machine translation of rare words with subword units,” 2016.
- [41] T. Wu, M. T. Ribeiro, J. Heer, and D. S. Weld, “Polyjuice: Generating counterfactuals for explaining, evaluating, and improving models,” 2021.
- [42] Z. Chen, J. Chen, Y. Chen, H. Yu, A. K. Singh, and M. Sra, “Lmexplainer: Grounding knowledge and explaining language models,” 2024.
- [43] G. Sarti, N. Feldhus, L. Sickert, and O. van der Wal, “Inseq: An interpretability toolkit for sequence generation models,” in *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 3: System Demonstrations)*, Association for Computational Linguistics, 2023.
- [44] D. N. Palacio, D. Khati, D. Rodriguez-Cardenas, A. Velasco, and D. Poshyvanyk, “On explaining (large) language models for code using global code-based explanations,” 2025.
- [45] E. Ameisen, J. Lindsey, A. Pearce, W. Gurnee, N. L. Turner, B. Chen, C. Citro, D. Abrahams, S. Carter, B. Hosmer, J. Marcus, M. Sklar, A. Templeton, T. Bricken, C. McDougall, H. Cunningham, T. Henighan, A. Jermyn, A. Jones, A. Persic, Z. Qi, T. B. Thompson, S. Zimmerman, K. Rivoire, T. Conerly, C. Olah, and J. Batson, “Circuit tracing: Revealing computational graphs in language models,” *Transformer Circuits Thread*, 2025.
- [46] O. Sanseviero and P. Schmid, “Introducing gemma 3: The developer guide,” 3 2025.
- [47] Q. Team, “Qwen2.5 technical report,” *arXiv preprint arXiv:2412.15115*, 2024.
- [48] Meta, “The llama 4 herd: The beginning of a new era of natively multimodal ai innovation,” 4 2025.